

Homework # 2

Mitchell Scott
(mtscot4)

1 Environment

1. List the environment you used to run the code (Operating System, CPU, RAM, etc.) .

Solution. The local environment where I was running the code was on a 2022 Mac-Book Pro with Apple M2 chip, 8 GB of RAM running macOS Sequoia 15.3.1.

2 Sequential Code

1. The results, which are averages of five runs where the fastest and slowest are removed, are seen in Tab. 1.

Trial	small.bin	medium.bin	large.bin	huge.bin
1	6.5834e-05 (high)	0.000828625 (low)	1.23898 (low)	6.73952
2	3.0208e-05	0.00109362 (high)	1.24350	6.69056 (low)
3	3.0875e-05	0.0009965	1.32917 (high)	6.79279
4	2.8542e-05 (low)	0.000983875	1.25815	6.71300
5	2.9375e-05	0.000959958	1.26216	6.82485 (high)
Avg	3.01527e-5	0.000980111	1.25460	6.748437

Table 1: Sequential Code for different file sizes. All are reported in seconds.

3 Parallel Code

3.1 Run Time

3.1.1 small.bin

The results for `small.bin` are given in table 2.

3.1.2 medium.bin

The results for `small.bin` are given in table 3.

Trial	1 thread (s)	2 thread (s)	4 thread (s)	8 thread (s)
1	1.1875e-4	1.60708e-4	0.00127112	0.001822 (low)
2	1.22291e-4	1.74000e-4	0.000373667	0.00222171
3	8.9958e-05 (low)	1.66537e-4	0.000272833 (low)	0.002149
4	1.5875e-4	1.42042e-4 (low)	0.00167750 (high)	0.00444679 (high)
5	1.67583e-4 (high)	1.82500e-4 (high)	0.000644041	0.00199421
Avg	1.33264e-4	1.67082e-4	7.62943e-4	0.00212164

Table 2: Time (s) to sort data in `small.bin` dataset for varying number of threads.

Trial	1 thread (s)	2 thread (s)	4 thread (s)	8 thread (s)
1	0.00123242 (high)	0.00364746 (high)	0.00293096 (high)	0.00612125
2	0.00110038	0.00269762	0.00197646	0.00327346 (low)
3	0.00122571	0.00240858	0.00184267 (low)	0.0141560 (high)
4	0.000960625 (low)	0.00208879 (low)	0.00202196	0.00400771
5	0.00100529	0.00277246	0.00209833	0.00410279
Avg	0.00111046	0.00262622	0.00203225	0.00474392

Table 3: Time (s) to sort data in `medium.bin` dataset for varying number of threads.

3.1.3 large.bin

The results for `large.bin` are given in table 4.

Trial	1 thread (s)	2 thread (s)	4 thread (s)	8 thread (s)
1	1.14397	0.908378 (low)	0.529193 (low)	0.534565(low)
2	1.13633	0.952551 (high)	0.534829	0.696034 (high)
3	1.17768 (high)	0.909059	0.532734	0.586798
4	1.1341	0.914919	0.533748	0.576026
5	1.13208 (low)	0.91012	0.543836 (high)	0.607107
Avg	1.13813	0.911366	0.533770	0.589977

Table 4: Time (s) to sort data in `large.bin` dataset for varying number of threads.

3.1.4 huge.bin

The results for `huge.bin` are given in table 5.

3.2 Speedup

In terms of the speedup, first we need to recall that

$$\text{Speedup} = \frac{T_{\text{sequential}}}{T_{\text{parallel}}} \quad (1)$$

So simply by dividing our times from the result in section 1, the speedup is seen in Tabs. 2, 3, 4 and 5, we arrive at the speedup, which is seen in Tab. 6.

This data is also presented in Fig. 1. As we can see from both the figure and the table, we have no great speedup for any number of processors with the `small` and the `medium`

Trial	1 thread (s)	2 thread (s)	4 thread (s)	8 thread (s)
1	5.99075 (high)	4.24292 (low)	2.60889 (low)	2.49775
2	5.96154	4.34822	2.76255 (high)	2.54355 (high)
3	5.88652 (low)	4.27647	2.61821	2.34084
4	5.94369	4.40685 (high)	2.69274	2.25988 (low)
5	5.98632	4.27752	2.71596	2.42781
Avg	5.96385	4.30074	2.67564	2.38880

Table 5: Time (s) to sort data in `huge.bin` dataset for varying number of threads.

# Threads	small	medium	large	huge
1	0.2263	0.8826	1.1023	1.1316
2	0.1805	0.3732	1.3766	1.5691
4	0.0395	0.4823	2.3504	2.5222
8	0.0142	0.2066	2.1265	2.8250

Table 6: Speedup of different data sets for varying number of threads.

dataset. This is due to the parallel overhead and message transferring, which isn't worth it at this size. However, for the **large** and the **huge** dataset, we see advantages for using the parallel code. Another note is the large dataset is only large enough for 4 processors to make sense to be used. It is large enough for some speedup but not too large to justify the extra communication costs that 8 have.

3.3 Data Splits

Ideally, we would split the data in such a way that all processors get the same amount of data to sort locally, but since we don't know the distribution of the data, we cannot *a priori* know the best splitting vector. Instead, we set a sampling parameter k , which for all the experiments was $k = 5$, which then is how many samples we took from each processor and sent back to the root processor. These samples were then concatenated and sorted to find the partition values. As we see from Fig. 2, this doesn't always work well as there are some processors that have more than their fair share of data to sort, while other processors are waiting around.

This leads to a nice extension that is: what is the optimal k value so that the workload is evenly balanced. Obviously there is a tradeoff between larger k meaning the workload is more balanced, and larger k means more computation.

4 Implementation

1. Explain the Algorithm you used.

Solution. First I used `MPI_Bcast` to broadcast the amount of data that will be sent to each processor dynamically. Once all processors had the amount that they would receive, I used `MPI_Scatterv` to actually send the data to each processor. The 'v' indicates that the data can be variable for each processor. Here we use `MPI_Reduce` to find the max and min of each processors' data. I took these values and the first k elements on each processor and used `MPI_Gather` to send these to the root processor,

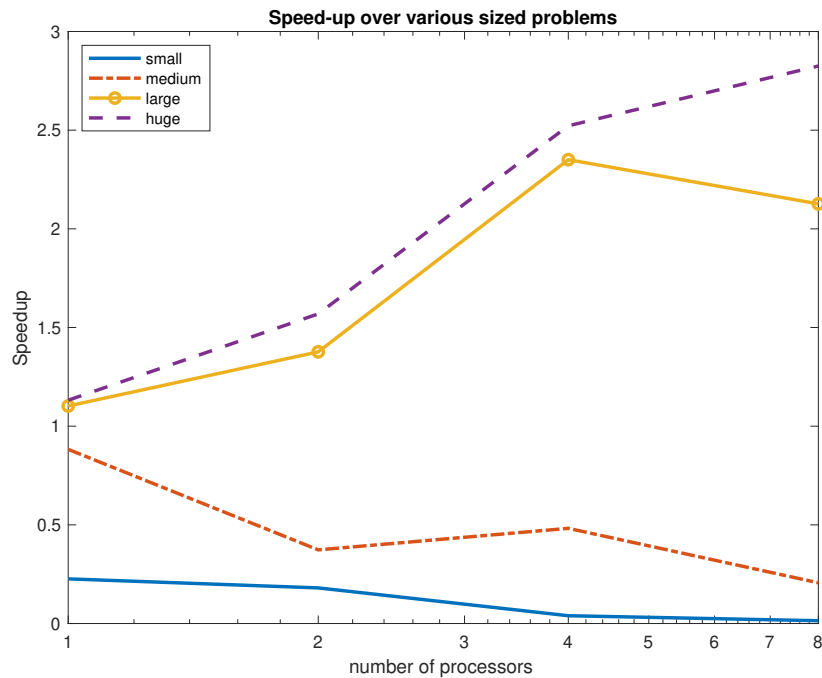


Figure 1: Comparing the speedup curves for different datasets.

where it can be sorted. Here we get the min and max of the data globally. Then from this sorted vector of samples, we use `MPI_Bcast` to send the splitting points back to each processor. We then figure out, based on the splitting vector, which data will be sent where and use `MPI_Alltoall` to send the send counts and receive counts of all the processes to each other process. We then use `MPI_Alltoallv` to actually send the data to the correct processor. This data is then sorted. Lastly, we use `MPI_Gather` and `MPI_Gatherv` to send the amount of data and the actual data to the root processor.

2. What is the performance bottleneck of your parallel implementation?

Solution. The major consideration here is the communication between different processors. While we can see this many places, the most prominent to me is the difference between 4 and 8 processors in the `large.bin` dataset. We actually notice a decrease in speed between these two cases, which is uncommon, since we have more threads doing the work. But by having double the communication costs, we see that this is stronger than the speedup by the extra threads.

Additionally a bottleneck is the load imbalance. As we see in the different data splits (Fig. 2), the data might be close to balanced but it is not exact. This allows some processors to do a lot of work while the others are doing nothing. This means the clock time is going to reflect the longer time it took that single processor to do all the work. This would be an interesting extension, to see what sampling parameter k is optimal for each dataset.




Figure 2: Visualized the data splits between processors for all datasets using (a) 2 threads, (b) 4 threads, and (c) 8 threads.

Acknowledgements

I would like to acknowledge that I worked with fellow CS 581 students Emma Hart and Francesco Brarda in OH as well as in the office on Thursday and Friday, February 27-28.